Regularize attention by how far it actually moves, not by its drop rate

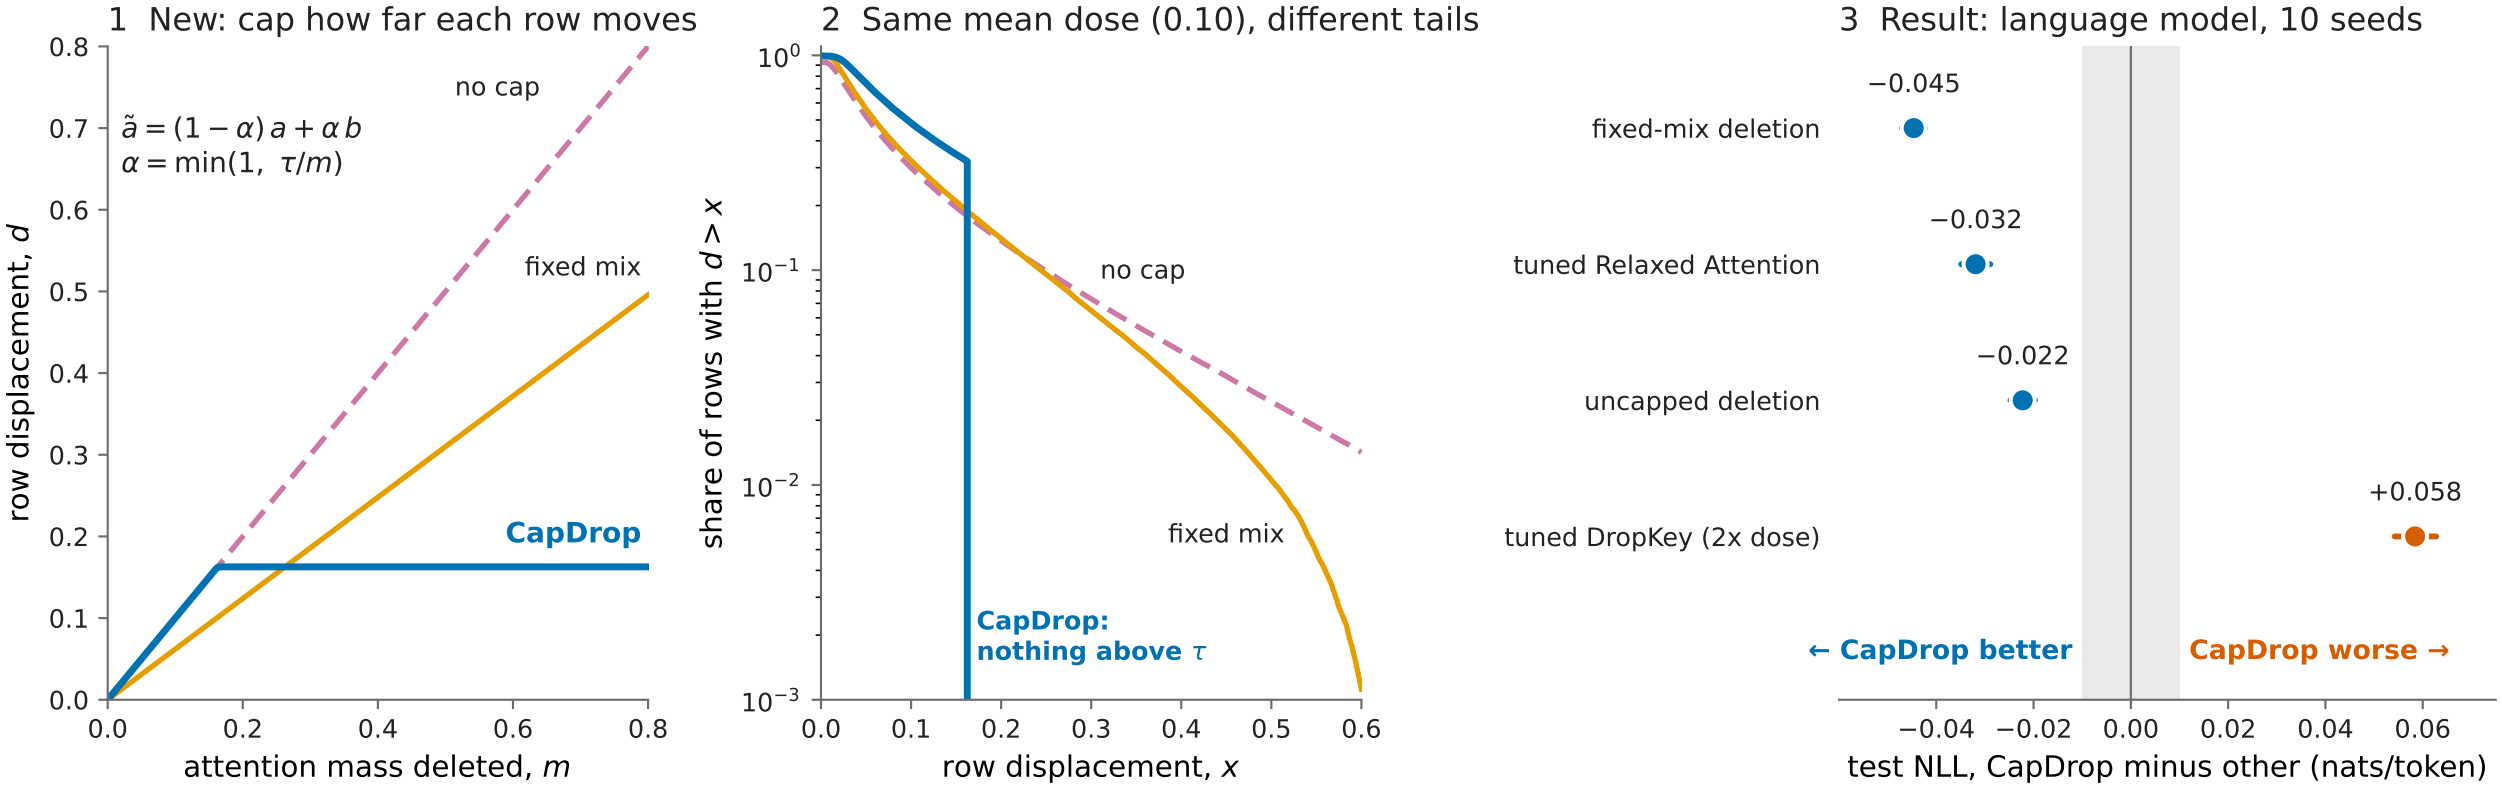

PubMedQA fine-tuning (BiomedBERT, 450 questions per fold): no attention regularizer, CapDrop included, improved on the unregularized model.
All arms matched on realized mean displacement; sealed test sets scored once; Holm-corrected exact tests; every run reported.